

Flush-Based Attacks


Today

- Flush+Reload
- Recovering missing bits from RSA keys
- Performance degradation attacks
- Flush+Flush

Bridging the gap

Cache utilises locality to bridge the gap

- Divides memory into *lines*
- Stores recently used lines
- In a *cache hit*, data is retrieved from the cache
- In a *cache miss*, data is retrieved from memory and inserted to the cache



Cache Consistency

- Memory and cache can be in inconsistent states
 - Rare, but possible
- Solution: Flushing the cache contents
 - Ensures that the next load is served from the memory

Processor





Cache




Memory

The FLUSH+RELOAD Attack


- Flush a memory line from the cache
- Wait a bit
- Measure time to load line
 - slow- \rightarrow no access
 - fast- \rightarrow access
- Repeat



Flush+Reload

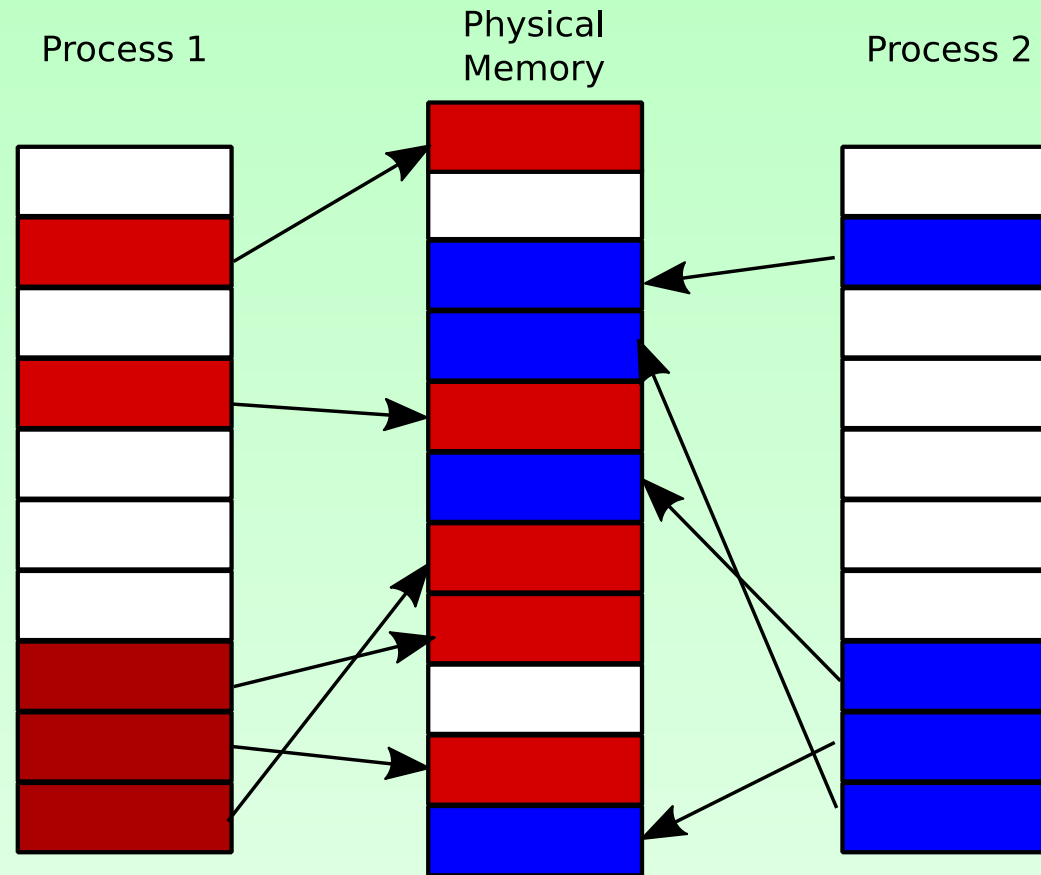


Flush+Reload



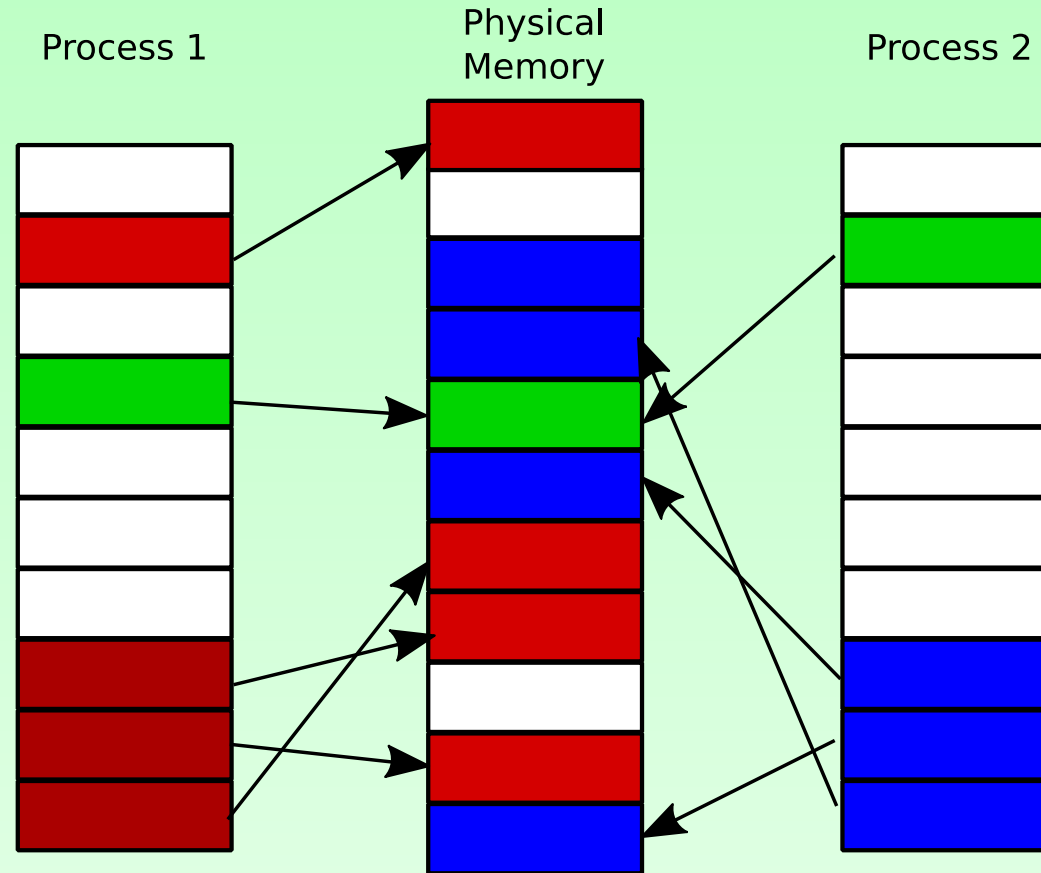
Detour - Virtual Memory

- Processes execute within a virtual address space
 - Virtual pages map to physical frames



Sharing

- Frames can be shared by multiple processes
 - Read only sharing maintains *functional* isolation
 - Protection using *Copy-on-write*



Causes of sharing

- Content-aware sharing
 - Pages from the same file have identical content
 - Shared program or library code
 - Can also share constant data
 - Shared images in PaaS clouds
- Content-based sharing (a.k.a. page deduplication)
 - The system identifies and coalesces identical pages
 - Implemented in many hypervisors and in most modern operating systems



Demo

The RSA Encryption System

- The RSA encryption is a public key cryptographic scheme



$$M = C^d \bmod N$$

M

$$C = M^e \bmod N$$



Key Generation:

- Select random primes p and q
- Calculate $N = pq$
- Select a public exponent $e (=65537)$
- Compute $d = e^{-1} \bmod \phi(N)$
- (N, e) is the public key
- (p, q, d) is the private key



Square and Multiply Exponentiation

```

 $x \leftarrow 1$ 
for  $i \leftarrow |d|-1$  downto 0 do
   $x \leftarrow x^2 \bmod n$ 
  if  $(d_i = 1)$  then
     $x = xC \bmod n$ 
  endif
done
return  $x$ 

```

Example:

$$11^5 \bmod 100 =$$

$$161,051 \bmod 100 = 51$$

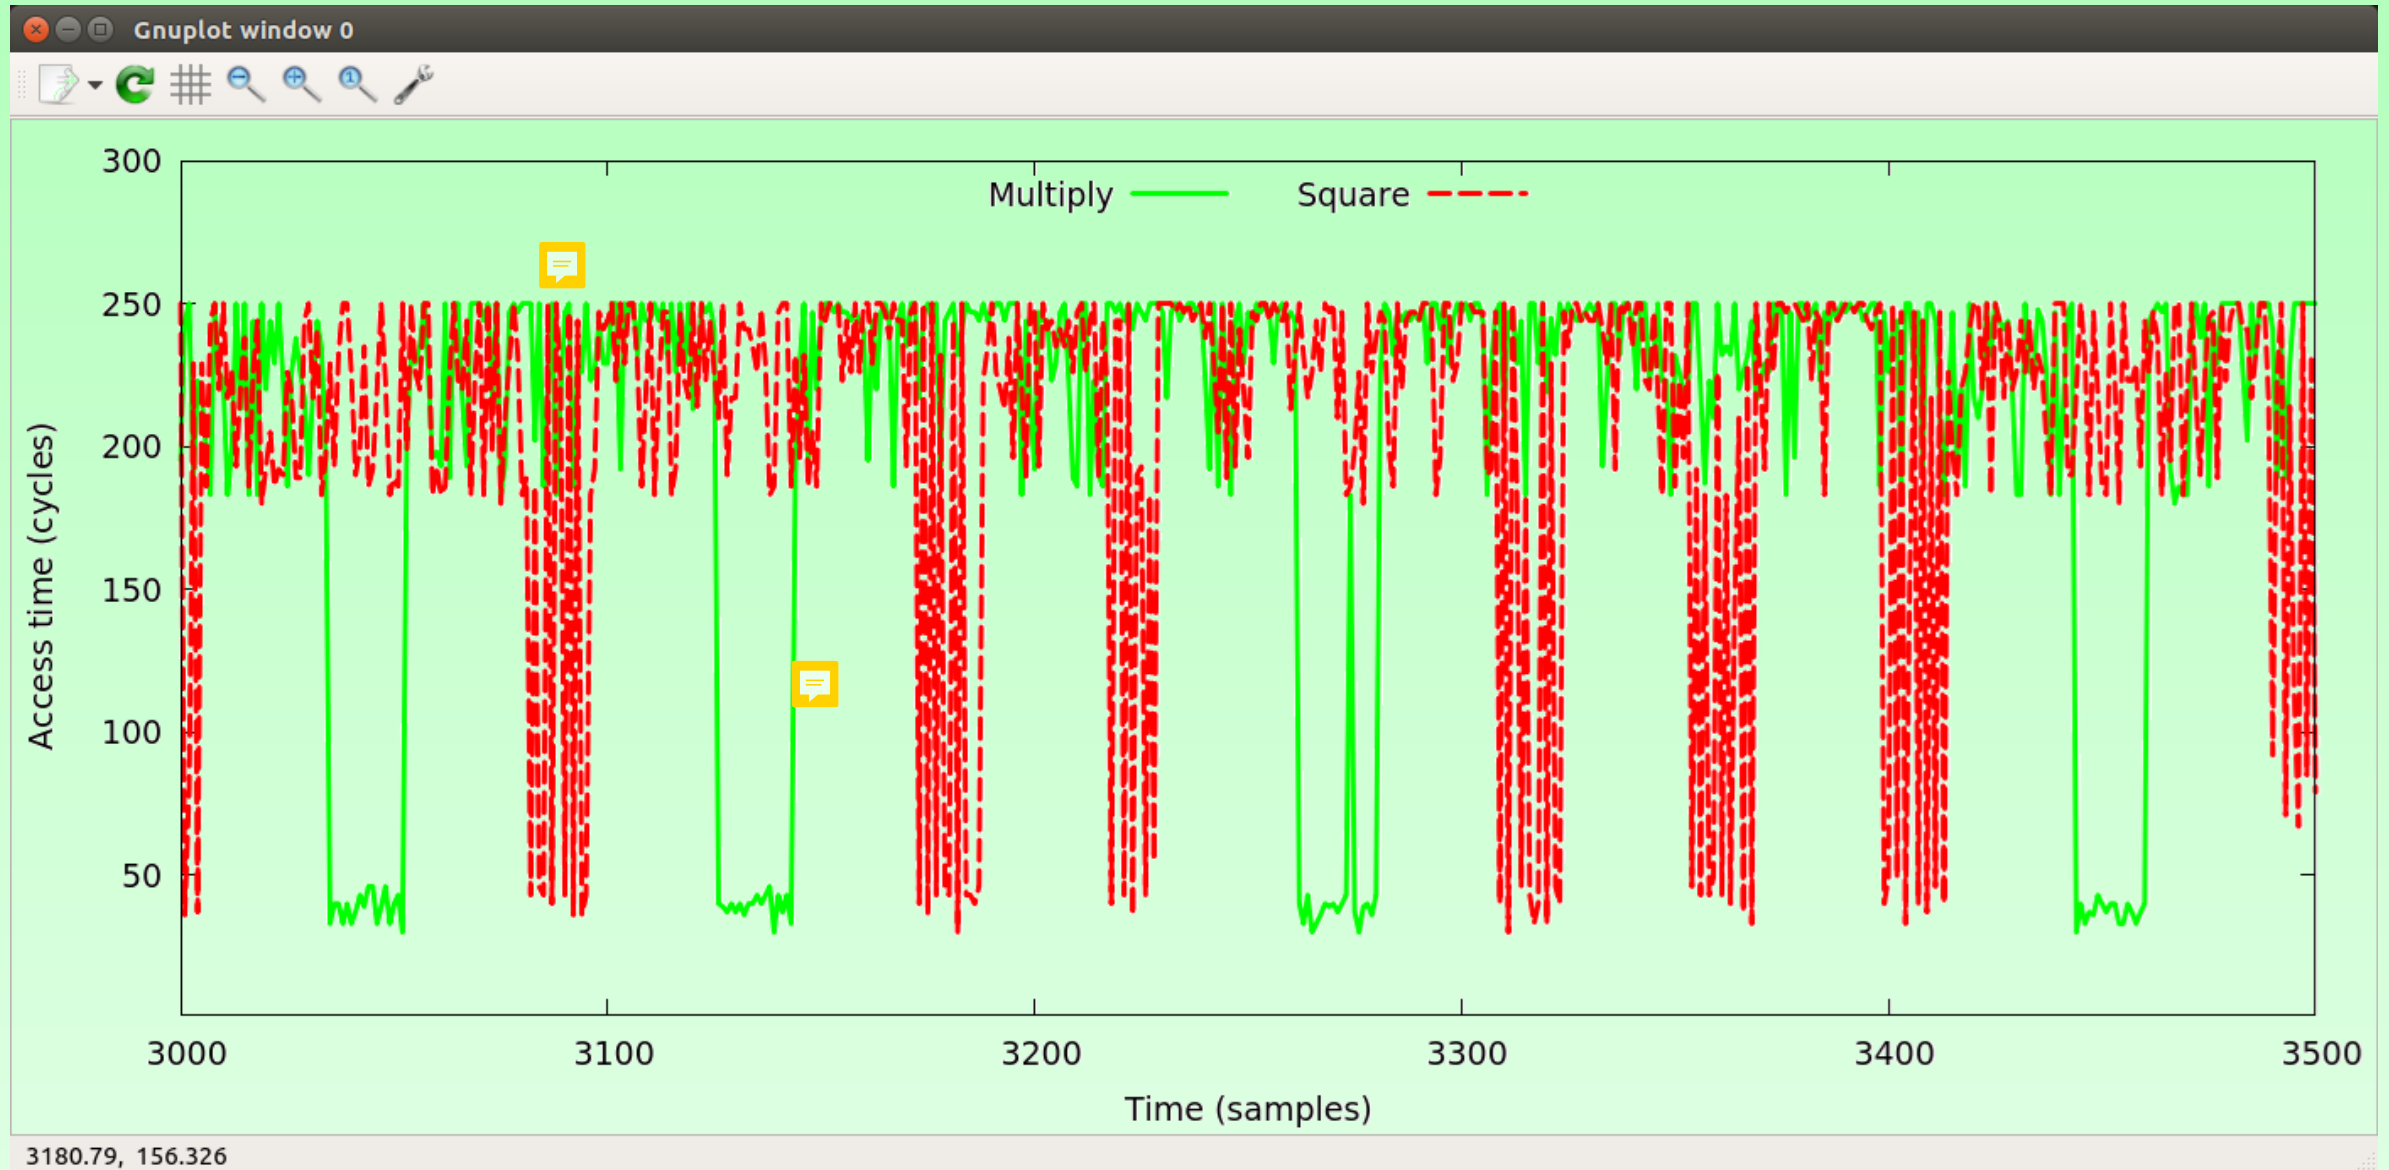
Operation	x	i	d_i

The private key is encoded in the sequence of operations !!!

**The private
key is
encoded in
the sequence
of operations
!!!**

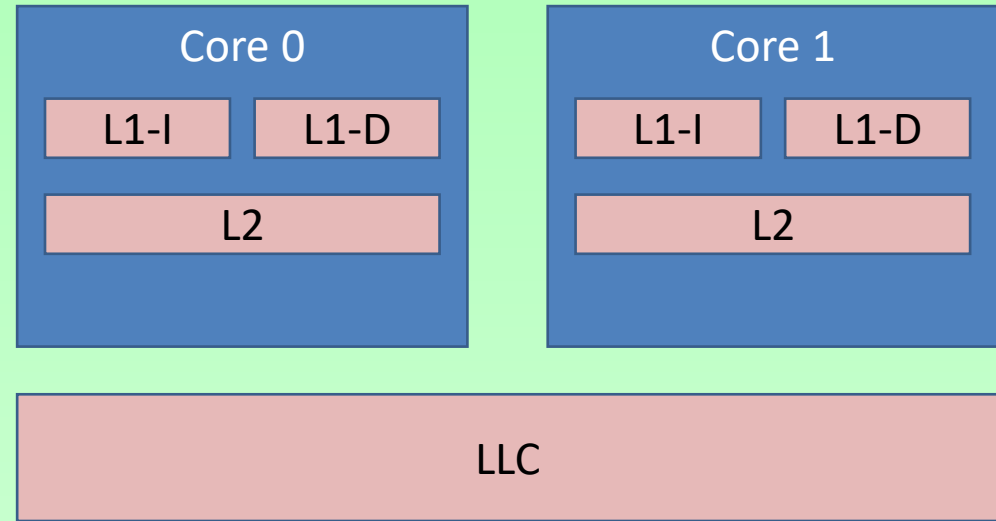
Demo

Flush+Reload on Square-and-Multiply



Why not AMD?

- Multilevel cache hierarchy
- Inclusive LLC: stores copies of all data in higher level caches
- Exclusive LLC: does not store data cached in higher level caches
- Non-inclusive LLC: not strictly inclusive. (Can be exclusive.)






Defences

Performance Degradation


Allan et al. “Amplifying Side Channel Through Performance Degradation”
ACSAC 2016

See also: Cabrera Aldaya and Brumley. “HyperDegrade: From GHz to MHz Effective CPU Frequencies”, USENIX Security 2022

Flush+Reload



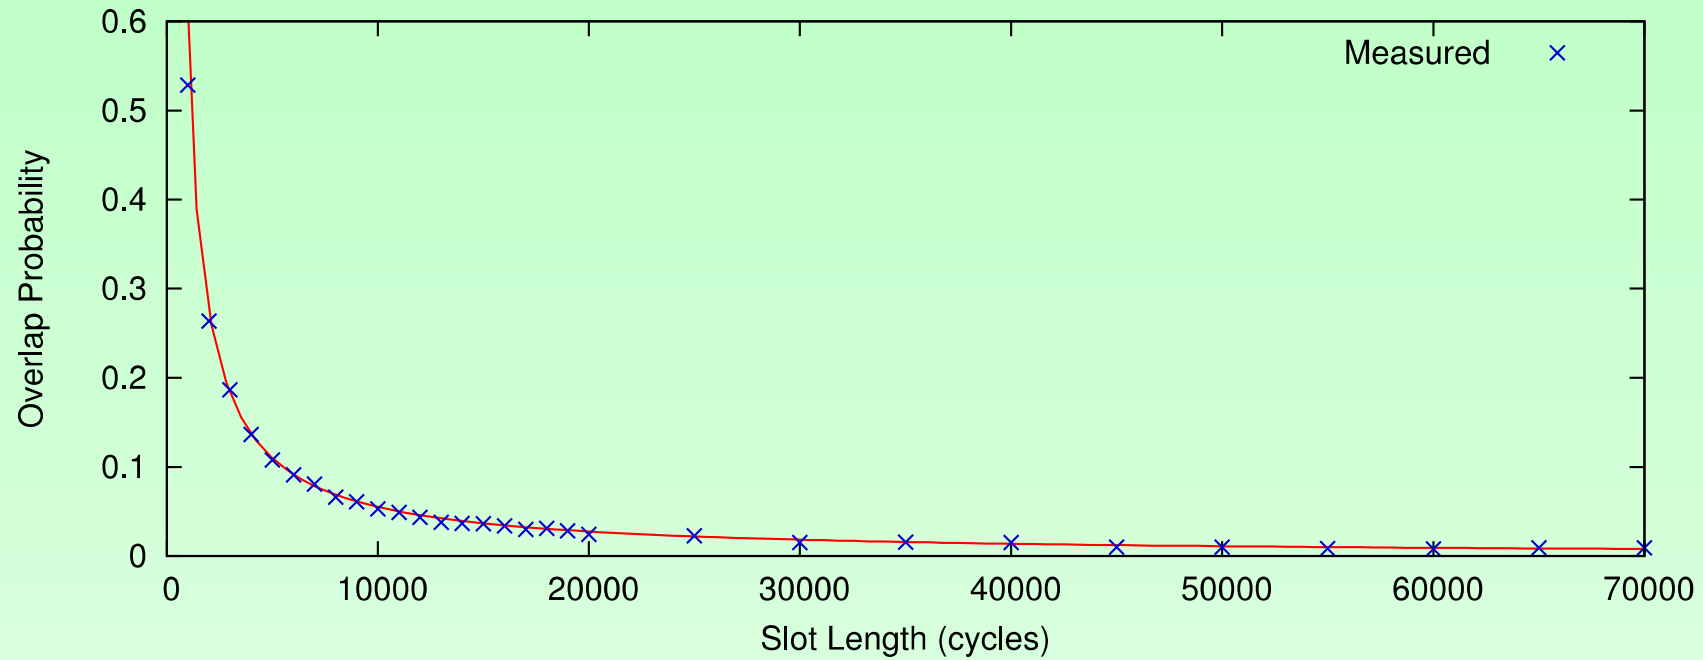
Timing matters



Access missed due to
temporal overlap

Probability of a probe miss

- Ratio of probe time to slot length




Probing OpenSSL ECDSA

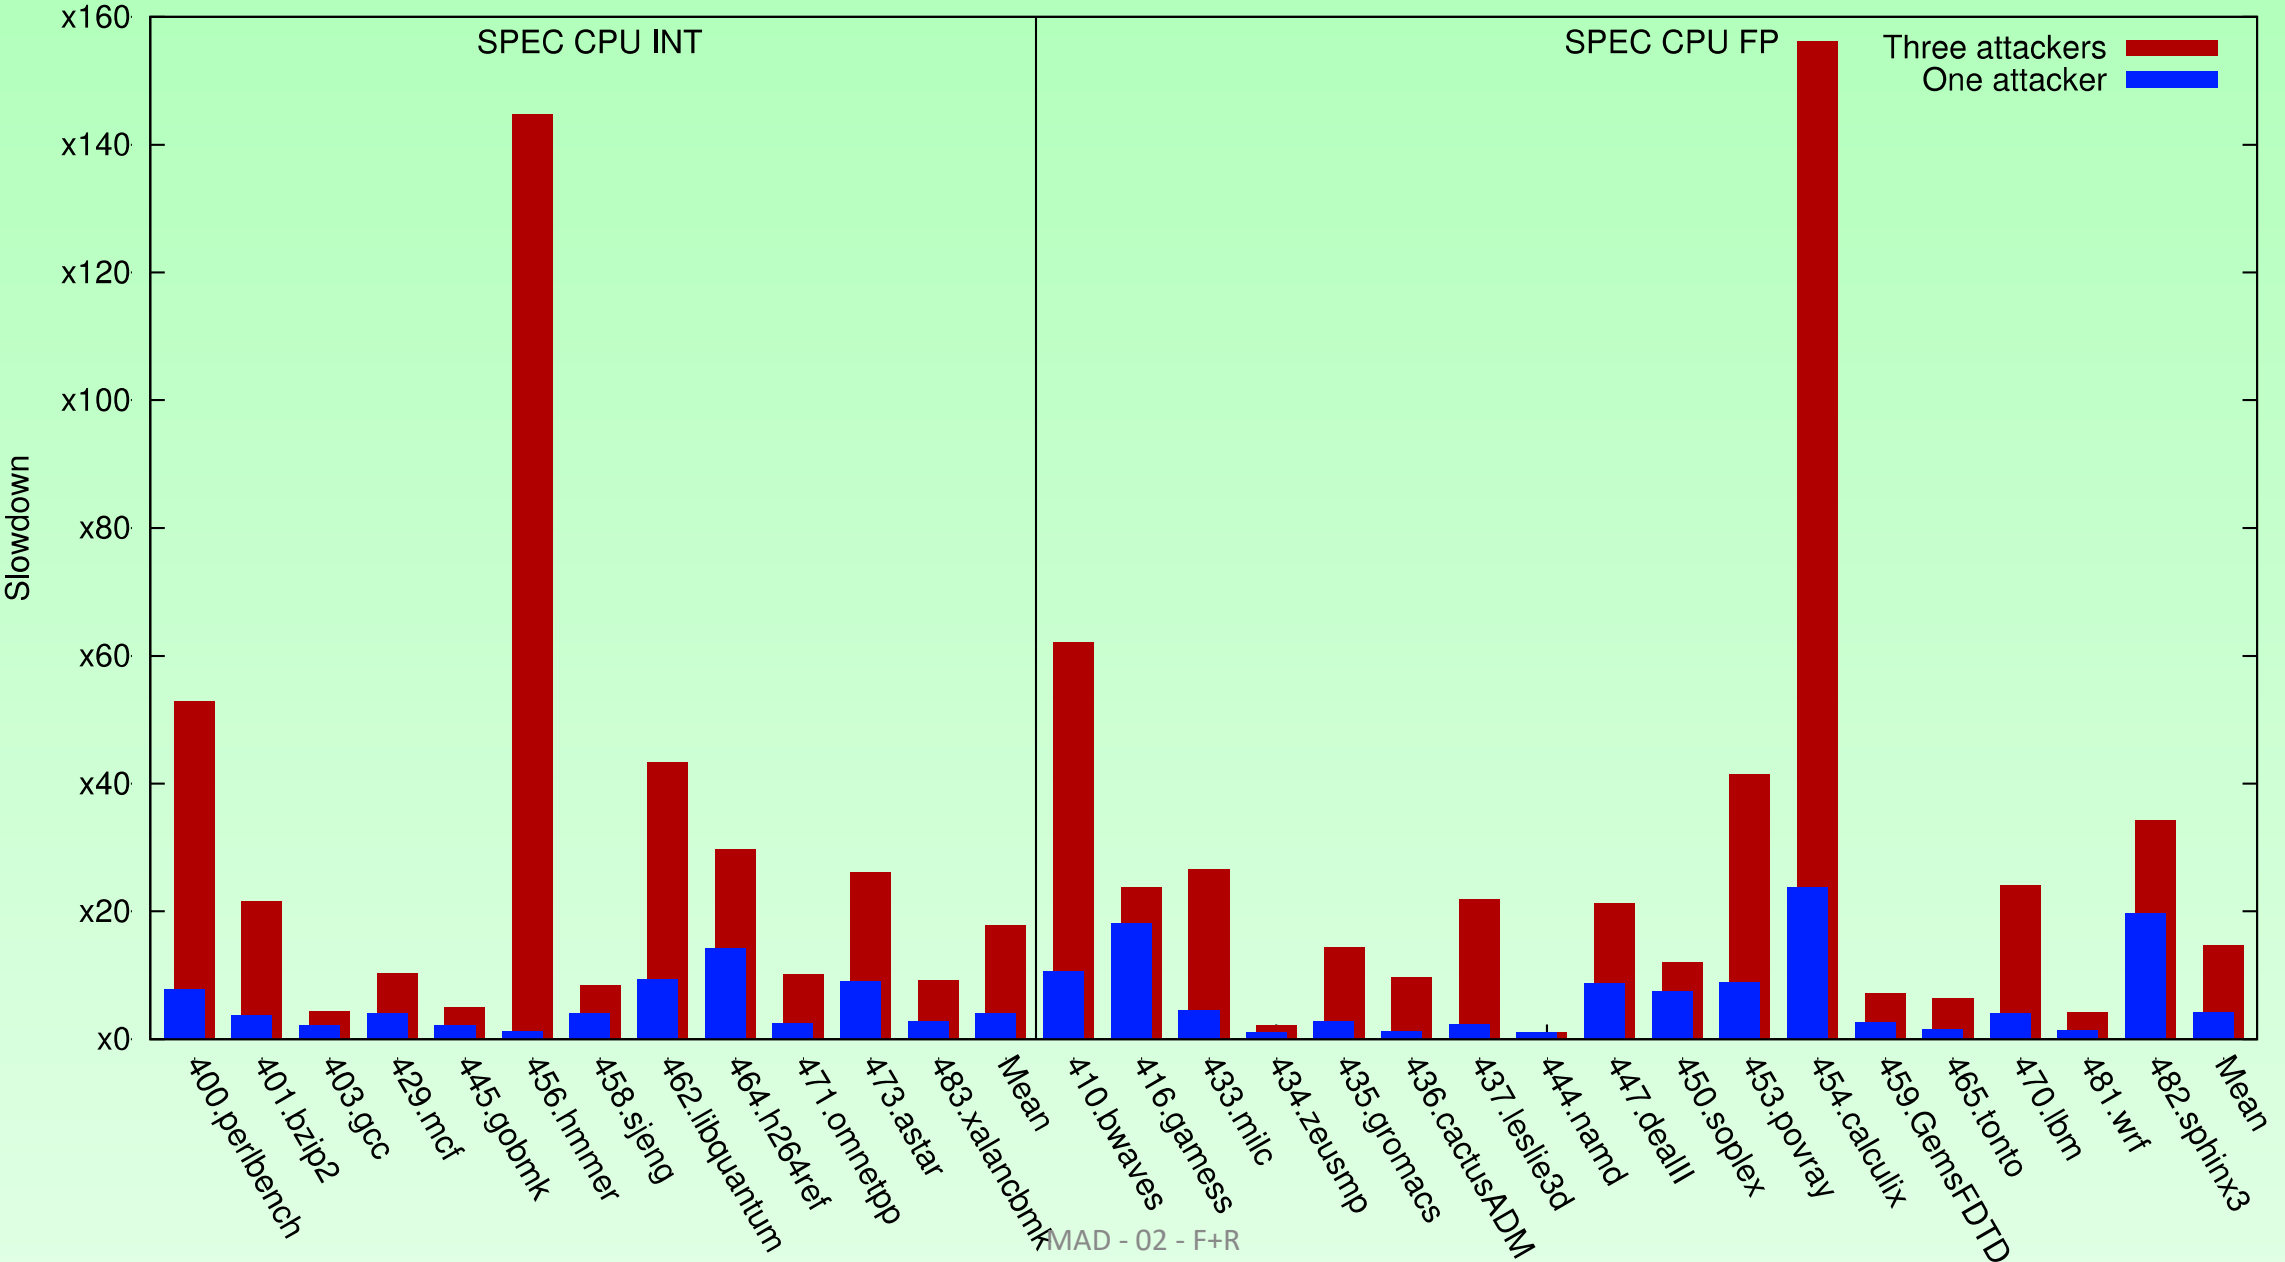
- For secp256k1, operation times are:
 - Addition: 2,892 cycles
 - Doubling: 3,086 cycles
 - Everything else: negligible
- To avoid aliasing errors, need to probe at least once every 1,450 cycles
- Limitations on probing
 - Each probe takes 450 cycles.
 - At the required probe rate, we can expect a 37% miss rate

Solution: slow the victim down

- Flushing data or code forces the victim to retrieve it from the memory
- Performance degradation attack:
 - Find frequently used code lines
 - For example, using profiling information
 - Repeatedly flush them from the cache



SPEC Results



Attacking OpenSSL


- Use two threads to flush three addresses spanning the inner loop in the bignum multiplication code
 - Add slowed by a factor of 53 (153,709 cycles)
 - Double slowed by a factor of 41 (126,282 cycles)
- Monitor four memory locations: one in add, one in double and two in the call chain of invert
- Probe each monitored location once every 17,000 cycles

Flush+Flush

Gruss et al. “Flush+Flush: A Fast and Stealthy Cache Attack”, DIMVA 2016

Flush+Flush

- Flush time depends on cache residency



- Small differences, variable baseline

Improvements

- Timing depends on core and slice. Can improve accuracy through calibration
 - Didier and Maurice, “Calibration Done Right: Noiseless Flush+Flush Attacks.”, DIMVA 2021
- Avoiding fences increases difference between hit and miss
 - Ding et al., “A Cross-Platform Cache Timing Attack Framework via Deep Learning”, DATE 2022

Summary

- [illegible]