

Heninger-Shacham

RSA

- Secret random primes p and q
- Calculate public $N = pq$
- For all r , $N \equiv pq \pmod{r}$
 - Specifically for all m , $N \equiv pq \pmod{2^m}$
- We know the LSB of p and q
 - Can we guess the next bit?
- What if we know the m LSBs? Can we guess the next bit?

Guessing bits



Let's try

$N =$ 1 0 0 1 0 1 0 1 1 0 1 0 1 1 1 0 1 0 1

$p =$ x x x x x x x x x x 1

x x x 1

But what if we
know th

- No known bits \Rightarrow Guess one to determine other
- One known \Rightarrow determine other
- Two known \Rightarrow rule out prior wrong guesses

$p =$ x x x x

$q =$ x x x x

x x x 1 1

x x x 1 1

$p =$ xxxxxxxx001

$q =$ xxxxxxxx101

~~$q =$ xxxxxxxx001~~

$q =$ xxxxxxxx111

~~$p =$ xxxxxxxx111~~

~~$q =$ xxxxxxxx011~~

2^n Guesses ☹️

RSA

- Random primes p and q
- Calculate $N = pq$
- Select a public exponent $e(=65537)$
- Compute $d=e^{-1} \bmod \varphi(N)$
- Encrypt: $C=M^e \bmod N$
- Decrypt: $M=C^d \bmod N$

Small $e \rightarrow$ fast computation

Large $d \rightarrow$ slow computation

CRT-RSA

- $d_p = d \bmod (p-1), \quad d_q = d \bmod (q-1)$
- $m_p = m^{d_p} \bmod p, \quad m_q = m^{d_q} \bmod q$
- $h = q^{-1}(m_p - m_q) \bmod p$
- $m = m_q + hq$

• Hence:

- $N = pq$
- $ed_p = k_p(p-1)+1$
- $ed_q = k_q(q-1)+1$

• Moreover, $0 < k_p, k_q < e$, and they are all related

Heninger-Shacham (CRYPTO 2009)

- A technique for finding the RSA-CRT key from partial information



Sliding Window Exponentiation

- Represent the exponent d in a convenient form:

$$d = \sum d_i 2^i \text{ where } d_i \text{ is either 0 or is odd } 0 < d_i < 2^w$$

- Precompute odd powers of the base b

$$b[i] = b^i \bmod p \text{ for odd } 0 < i < 2^w$$

$$B[1] = b$$

$$b_{sqr} = b^2 \bmod p$$

for $i = 3, 5, \dots, 2^w - 1$ **do**

$$b[i] = b[i-2] \cdot b_{sqr} \bmod p$$

Sliding Window Exponentiation

- Perform the exponentiation

```
 $r \leftarrow 1$   
for  $i = |d|-1, \dots, 0$  do  
     $r \leftarrow r^2 \bmod p$   
    if  $d_i \neq 0$  then  
         $r \leftarrow r \cdot b[d_i] \bmod p$   
return  $r$ 
```

Sliding window representation revisited

$$d = \sum d_i 2^i \text{ where } d_i \text{ is either 0 or 1 and } 0 < d_i < 2^w$$

Another way of looking at this:

- Divide d into *windows*
 - Windows are at most w bits wide
 - Windows start and end with 1

$$\begin{aligned} d = 4312 &= \boxed{1} \ 0 \ 0 \ 0 \ 0 \ 0 \ \boxed{1} \ \boxed{1 \ 0 \ 1 \ 1} \ 0 \ 0 \ 0 \\ d_i &= \quad \quad \quad 1 \ 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \ 1 \ 1 \ 0 \ 0 \ 0 \\ \sum d_i 2^i &= 1 \cdot 2^{12} + 1 \cdot 2^7 + 11 \cdot 2^3 = 4096 + 128 + 11 \cdot 8 = 4312 \end{aligned}$$

Sliding window representation revisited

$$d = \sum d_i 2^i \text{ where } d_i \text{ is either 0 or 1, } 0 \leq d_i < 2^w$$


Not a unique representation

$$d = 4312 = \begin{array}{cccccccccccc} \boxed{1} & 0 & 0 & 0 & 0 & \boxed{1} & \boxed{1} & 0 & 1 & 1 & 0 & 0 & 0 \\ \boxed{1} & 0 & 0 & 0 & 0 & \boxed{1} & \boxed{1} & 0 & 1 & \boxed{1} & 0 & 0 & 0 \\ \boxed{1} & 0 & 0 & 0 & 0 & \boxed{1} & \boxed{1} & 0 & 1 & \boxed{1} & 0 & 0 & 0 \end{array}$$

Minimise the number of windows for best performance

A greedy algorithm

- Scan the bits of d until finding a 1
 - *Open* a window of length w bits
 - *Close* the window at the last 1
 - Repeat
-
- The greedy algorithm is optimal



Analysis

- Successive open positions are at least w bits apart

- Assuming a random d

1 0 0 0 0 1 1 0 1 1 0 0 0

- w bits apart with probability $1/2$
 - $w+1$ bits apart with probability $1/4$
 - $w+2$ bits apart with probability $1/8$
 - Etc...
- On average, successive open positions are $w+1$ bits apart
 - Expected number of windows is $|d|/(w+1)$


Left-to-right vs. right-to-left

- Previously, we ran the greedy algorithm from the right to the left.

1	0	0	0	0	1	1	0	1	1	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---

- Can also run from the left to the right

1	0	0	0	0	1	1	0	1	1	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---



- The analysis still applies
 - We get the same number of windows
- Can combine left-to-right with exponentiation

Recovering the exponent

- We know the positions of the windows

x x x x x x x x x x x x x

- Windows start with 1

1 x x x x 1 x x x 1 x x x

- Everything outside maximum window must be 0

1 0 x x x 1 x x x 1 0 0 0

- On average we get 2 bits per window or $2|d|/(w+1)$ per exponentiation

Heninger-Shacham and the side-channel results

- A technique for finding the RSA-CRT key from partial information
- On average, needs one bit of p , q , d_p , d_q
- In our case we do not know bits of p or q
- For a successful attack we need to know half the bits of d_p and d_q
- We get 2 bits per window of size $w+1$
 - For $w>3$, we need more information

Left-to-right leaks more

- Right to left ($d=8625$):

<u>1</u>	0	0	0	0	<u>1</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	0	0	0	<u>1</u>
<u>x</u>	x	x	x	x	<u>x</u>	x	x	x	<u>x</u>	x	x	x	<u>x</u>
<u>1</u>	0	x	x	x	<u>1</u>	x	x	x	<u>1</u>	x	x	x	<u>1</u>


- Left-to-right ($d=8625$):

<u>1</u>	0	0	0	0	<u>1</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	0	0	0	<u>1</u>
<u>x</u>	x	x	x	x	x	x	x	<u>x</u>	<u>x</u>	x	x	x	<u>x</u>
<u>1</u>	0	0	0	0	<u>x</u>	<u>x</u>	<u>x</u>	<u>1</u>	<u>1</u>	x	x	x	<u>1</u>
<u>1</u>	0	0	0	0	<u>1</u>	x	x	<u>1</u>	<u>1</u>	x	x	x	<u>1</u>
<u>1</u>	0	0	0	0	<u>1</u>	x	x	<u>1</u>	<u>1</u>	0	0	0	<u>1</u>

Results

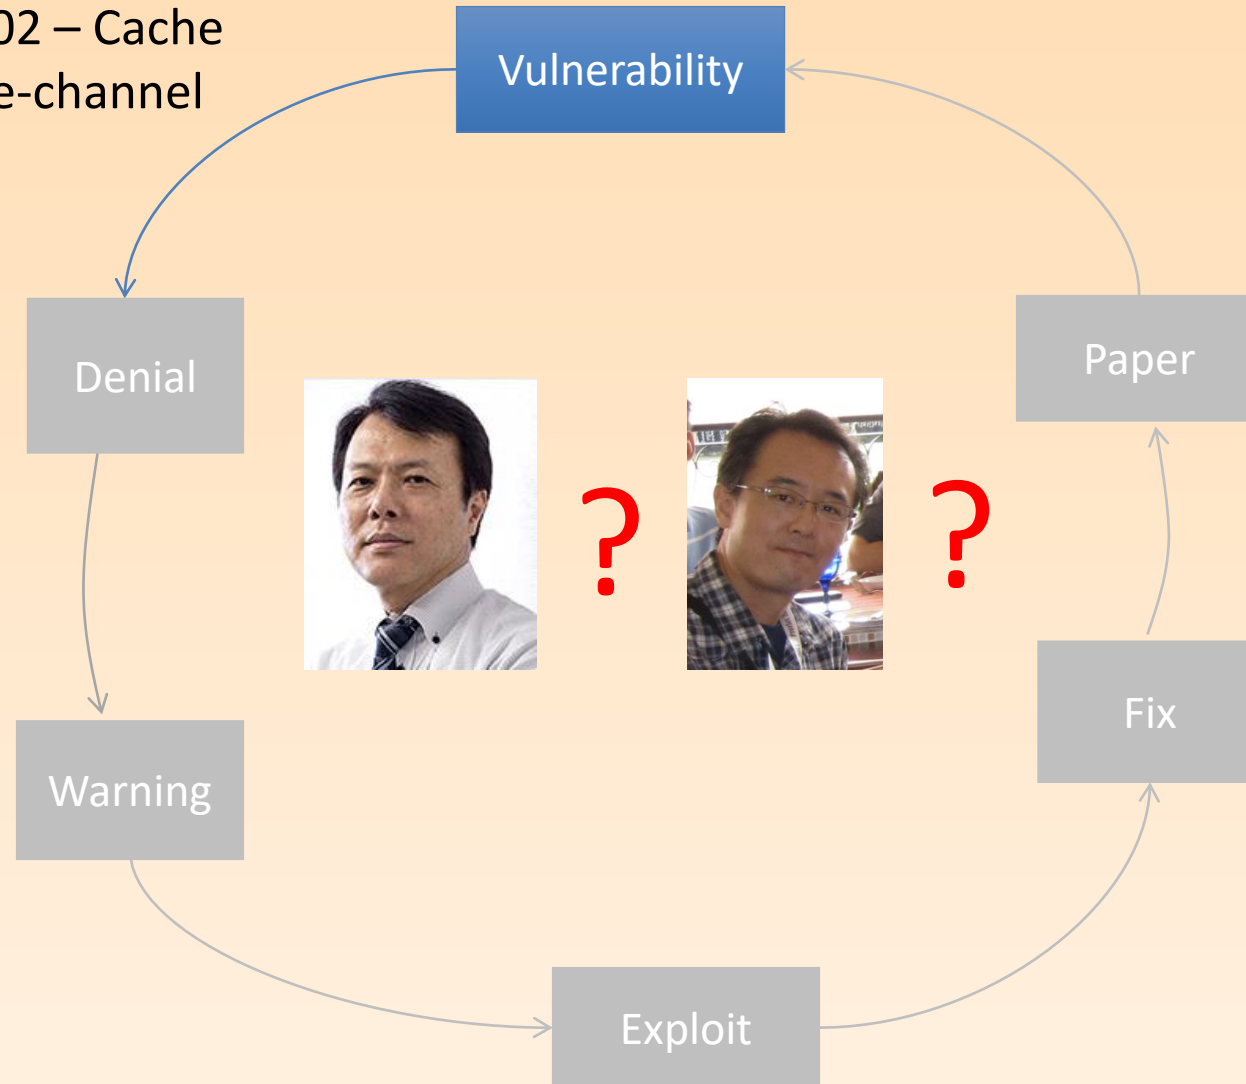
- For $w=4$ can get 2.725 bits per window
 - Break all 1024-bit RSA keys in Libgcrypt
- For $w=5$ can get 2.766 bits per window
 - Break 13% of the 2048-bit RSA keys
- For more information see
Bernstein et al. “Sliding Right Into Disaster: Left-to-Right Sliding Windows Leak”, CHES 2017

Attack Life Cycle




Round 1

Tsunoo, Tsujihara,
Minematsu and
Miyauchi 2002 – Cache
may leak side-channel
information





Round 1

Tsunoo, Tsujihara,
Minematsu and
Miyauchi 2002 – Cache
may leak side-channel
information



The RSA Encryption System

- The RSA encryption is a public key cryptographic scheme



Key Generation:

- Select random primes p and q
- Calculate $N = pq$
- Select a public exponent $e (=65537)$
- Compute $d = e^{-1} \bmod \phi(N)$
- (N, e) is the public key
- (p, q, d) is the private key

Fixed Window Exponentiation

- Divide exponent into *windows* of size w

1 0	0 1 1 0	1 0 0 1	0 1 0 0
-----	---------	---------	---------

$d_3=2$ $d_2=6$ $d_1=9$ $d_0=4$

- Precompute $b_i = b^i \bmod m$:

$$b_0 \leftarrow 1$$

$$b_1 \leftarrow b$$

for $i = 2, 3, \dots, 2^w - 1$ **do**

$$b_i \leftarrow b_{i-1} \cdot b \bmod m$$



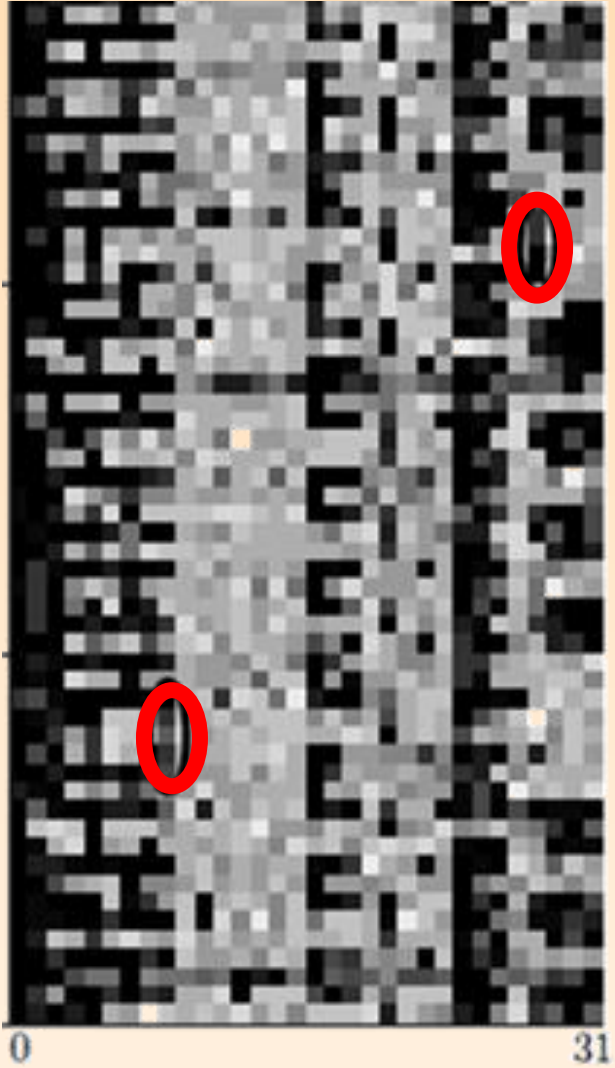
Calculating the exponent

```
 $r \leftarrow 1$   
for  $i = \lceil n/w \rceil - 1, \dots, 0$  do  
    for  $j = 1, \dots, w$  do  
         $r \leftarrow r \cdot r \bmod m$   
         $r \leftarrow r \cdot b_{d_i} \bmod m$   
return  $r$ 
```

1	0	0	1	1	0	1	0	0	1	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---

$d_3=2$ $d_2=6$ $d_1=9$ $d_0=4$

Prime+Probe against RSA (Percival 2005)



Memory

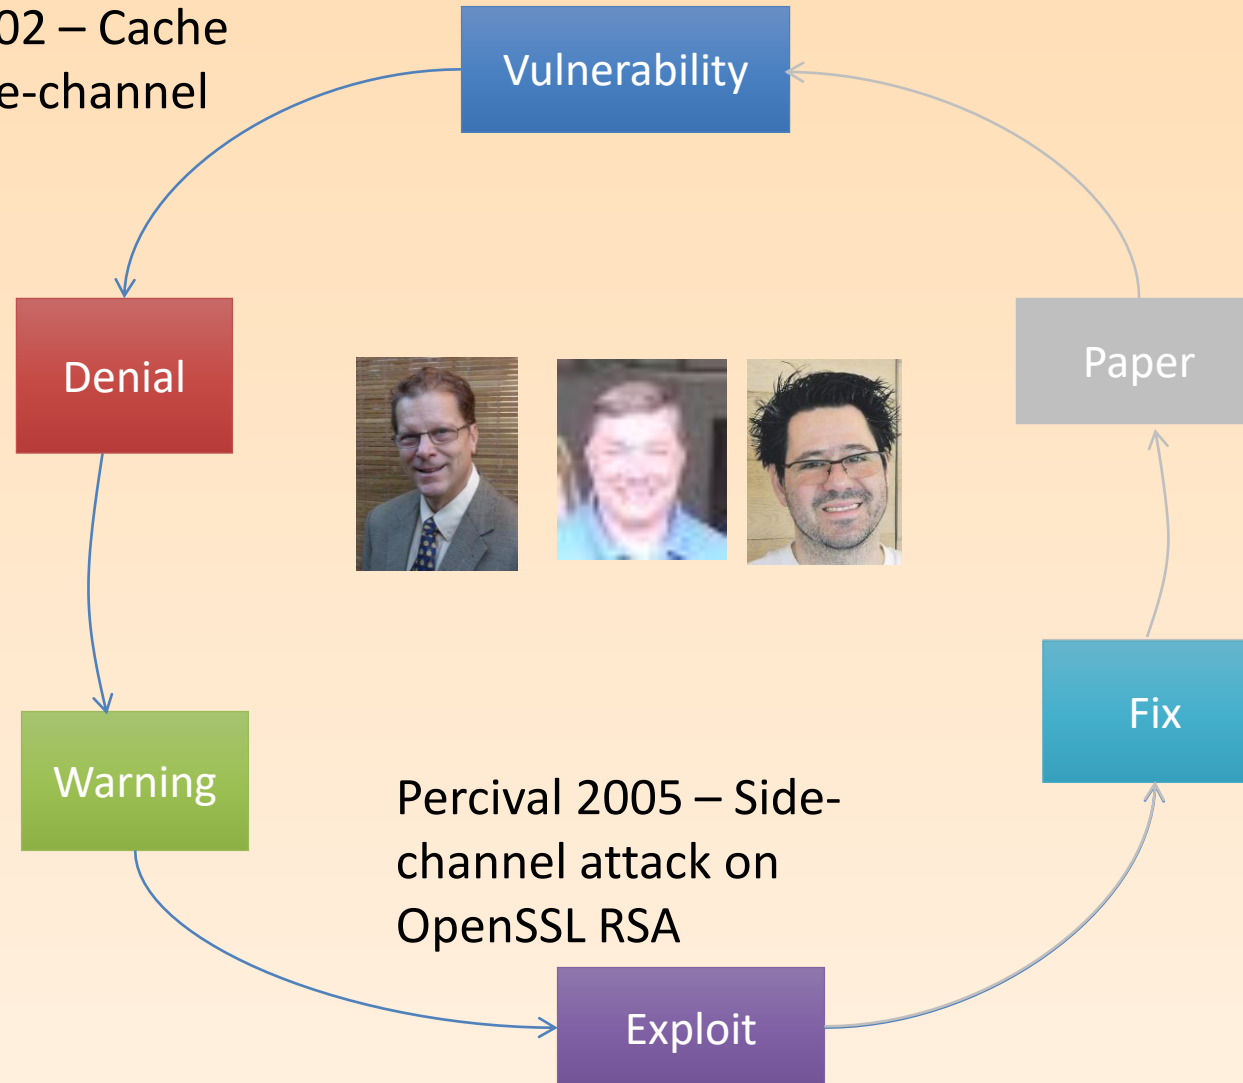
Why we can identify multipliers

- Each multiplier occupies consecutive cache lines
- Accessed throughout the multiplication

offset	0	1	2	...	63
Line 0	M ₀ [0]	M ₀ [1]	M ₀ [2]	...	M ₀ [63]
Line 1	M ₀ [64]	M ₀ [65]	M ₀ [66]	...	M ₀ [127]
Line 2	M ₀ [128]	M ₀ [129]	M ₀ [130]	...	M ₀ [191]
Line 3	M ₁ [0]	M ₁ [1]	M ₁ [2]	...	M ₁ [63]
Line 4	M ₁ [64]	M ₁ [65]	M ₁ [66]	...	M ₁ [127]
Line 5	M ₁ [128]	M ₁ [129]	M ₁ [130]	...	M ₁ [191]
Line 6	M ₂ [0]	M ₂ [1]	M ₂ [2]	...	M ₂ [63]
Line 7	M ₂ [64]	M ₂ [65]	M ₂ [66]	...	M ₂ [127]
Line 8	M ₂ [128]	M ₂ [129]	M ₂ [130]	...	M ₂ [191]
	•	•	•		•
	•	•	•		•
	•	•	•		•
Line 191	M ₆₃ [128]	M ₆₃ [129]	M ₆₃ [130]	...	M ₆₃ [191]

Round 1

Tsunoo, Tsujihara,
Minematsu and
Miyauchi 2002 – Cache
may leak side-channel
information



Brickell, Graunke
& Seifert 2006 –
Use Scatter
Gather

Percival 2005 – Side-
channel attack on
OpenSSL RSA

Scatter-Gather

- Mitigate Prime+Probe
 - Sequence of accesses to cache lines does not depend on secret data

offset	0	1	2	...	63	offset	0	1	2	...	63
Line 0	M ₀ [0]	M ₀ [1]	M ₀ [2]	...	M ₀ [63]	Line 0	M ₀ [0]	M ₁ [0]	M ₂ [0]	...	M ₆₃ [0]
Line 1	M ₀ [64]	M ₀ [65]	M ₀ [66]	...	M ₀ [127]	Line 1	M ₀ [1]	M ₁ [1]	M ₂ [1]	...	M ₆₃ [1]
Line 2	M ₀ [128]	M ₀ [129]	M ₀ [130]	...	M ₀ [191]	Line 2	M ₀ [2]	M ₁ [2]	M ₂ [2]	...	M ₆₃ [2]
Line 3	M ₁ [0]	M ₁ [1]	M ₁ [2]	...	M ₁ [63]	Line 3	M ₀ [3]	M ₁ [3]	M ₂ [3]	...	M ₆₃ [3]
Line 4	M ₁ [64]	M ₁ [65]	M ₁ [66]	...	M ₁ [127]	Line 4	M ₀ [4]	M ₁ [4]	M ₂ [4]	...	M ₆₃ [4]
Line 5	M ₁ [128]	M ₁ [129]	M ₁ [130]	...	M ₁ [191]	Line 5	M ₀ [5]	M ₁ [5]	M ₂ [5]	...	M ₆₃ [5]
Line 6	M ₂ [0]	M ₂ [1]	M ₂ [2]	...	M ₂ [63]	Line 6	M ₀ [6]	M ₁ [6]	M ₂ [6]	...	M ₆₃ [6]
Line 7	M ₂ [64]	M ₂ [65]	M ₂ [66]	...	M ₂ [127]	Line 7	M ₀ [7]	M ₁ [7]	M ₂ [7]	...	M ₆₃ [7]
Line 8	M ₂ [128]	M ₂ [129]	M ₂ [130]	...	M ₂ [191]	Line 8	M ₀ [8]	M ₁ [8]	M ₂ [8]	...	M ₆₃ [8]
	•	•	•		•		•	•	•		•
	•	•	•		•		•	•	•		•
	•	•	•		•		•	•	•		•
Line 191	M ₆₃ [128]	M ₆₃ [129]	M ₆₃ [130]	...	M ₆₃ [191]	Line 191	M ₀ [191]	M ₁ [191]	M ₂ [191]	...	M ₆₃ [191]

OpenSSL's Layout

offset	0	Bin 0	7 8	Bin 1	15	...	56	Bin 7	63
Line 0		M ₀ [0-7]		M ₁ [0-7]		...		M ₇ [0-7]	
Line 1		M ₈ [0-7]		M ₉ [0-7]		...		M ₁₅ [0-7]	
Line 2		M ₁₆ [0-7]		M ₁₇ [0-7]		...		M ₂₃ [0-7]	
Line 3		M ₂₄ [0-7]		M ₂₅ [0-7]		...		M ₃₁ [0-7]	
Line 4		M ₀ [8-15]		M ₁ [8-15]		...		M ₇ [8-15]	
Line 5		M ₈ [8-15]		M ₉ [8-15]		...		M ₁₅ [8-15]	
Line 6		M ₁₆ [8-15]		M ₁₇ [8-15]		...		M ₂₃ [8-15]	
Line 7		M ₂₄ [8-15]		M ₂₅ [8-15]		...		M ₃₁ [8-15]	
Line 8		M ₀ [16-23]		M ₁ [16-23]		...		M ₇ [16-23]	
		•		•				•	
		•		•				•	
		•		•				•	
Line 95		M ₂₄ [184-191]		M ₂₅ [184-191]		...		M ₃₁ [184-191]	

Round 1

Tsuruoka, a,
Mir
Miy
ma
info



Cache-timing attacks on AES

Daniel J. Bernstein *

Department of Mathematics, Statistics, and Computer Science (M/C 249)
The University of Illinois at Chicago
Chicago, IL 60607-7045
djb@cr.yp.to

Abstract. This paper demonstrates complete AES key recovery from known-plaintext timings of a network server on another computer. This attack should be blamed on the AES design, not on the particular AES

Denial

Bernstein 2005, Osvik
Shamir & Tromer 2006
– Scatter Gather may
leak information

Paper

Cache Attacks and Countermeasures: the Case of AES (Extended Version) revised 2005-11-20

Dag Arne Osvik¹, Adi Shamir² and Eran Tromer²



act.
te of
for cryptan
allow an un
despite par

le-cha
leakage
s that
her pr
very p

ess leakage through
s, which can be used
lookups. The attacks
the same processor,
ualization. Some of

er attack on
OpenSSL RSA


Exploit

Fix

Brickell, Graunke
& Seifert 2006 –
Use Scatter
Gather


Cache banks

- To support superscalar processing the cache is divided into cache banks
 - Bits 2-5 of the address determine the bank

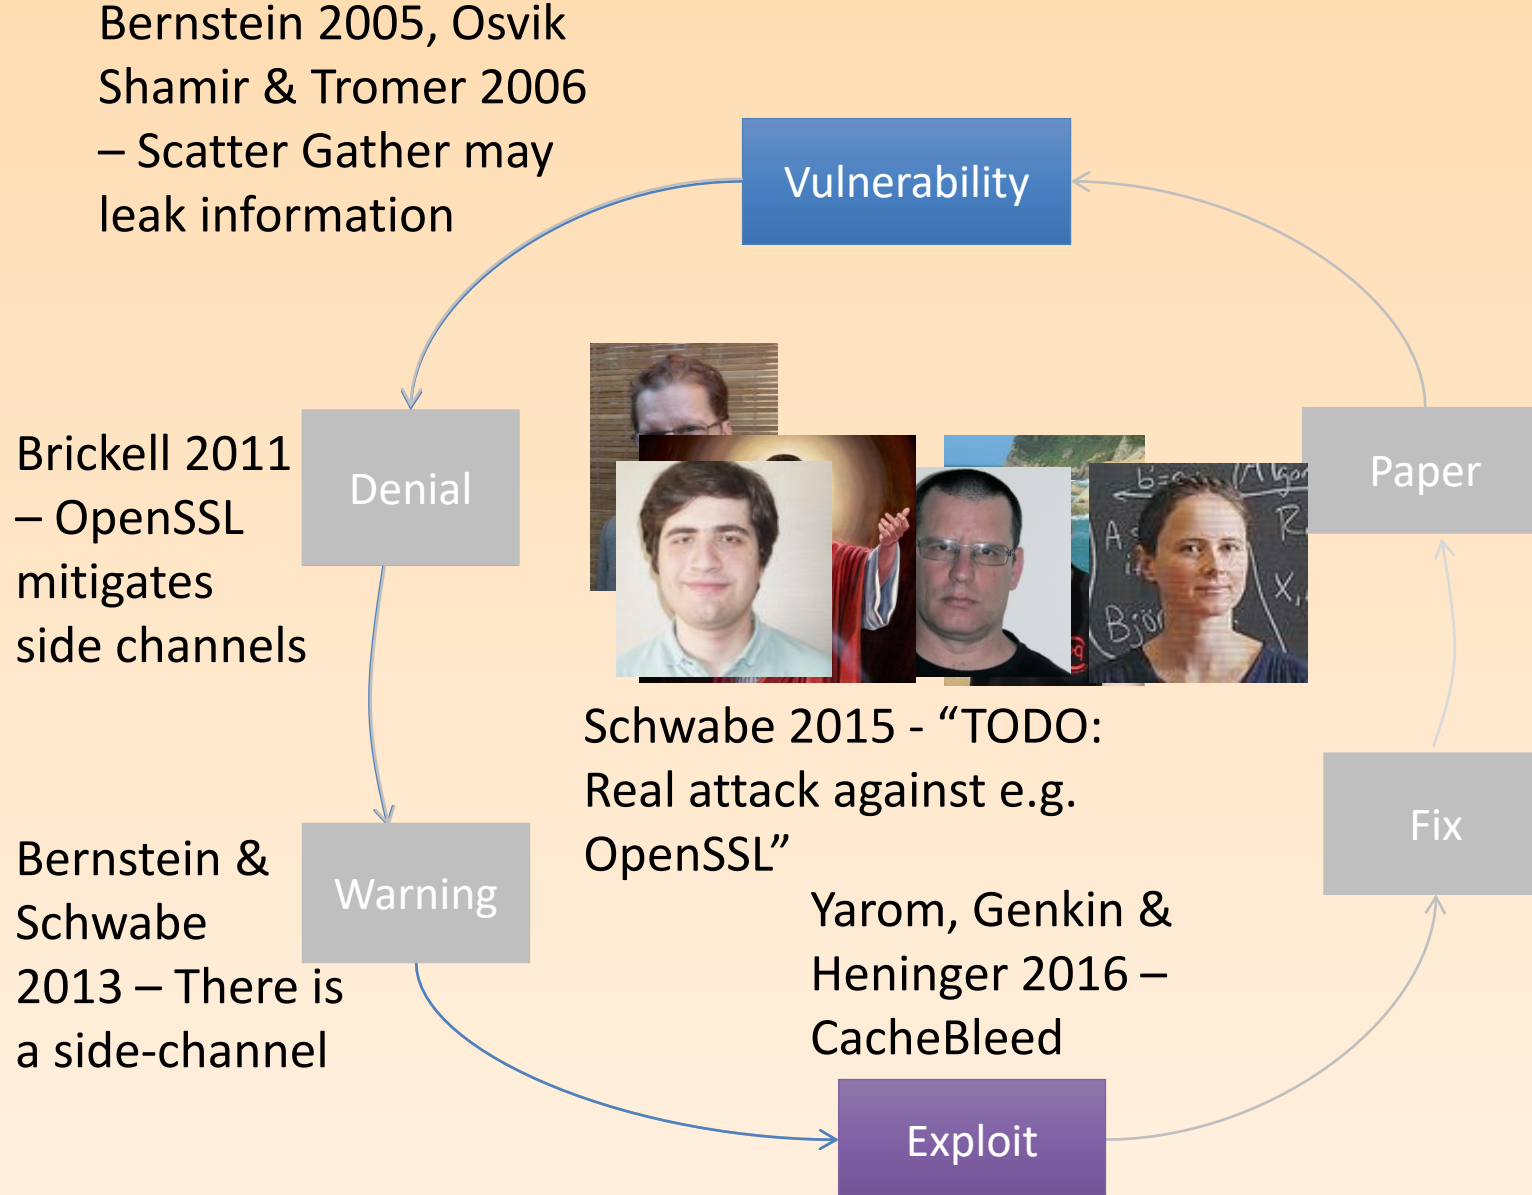


Cache Banks

- In Sandy Bridge, each bank can serve only one request per cycle.
 - Concurrent access to different banks is always possible
 - Concurrent access to the same bank causes delays




Round 2




CacheBleed Operation

- Spy generate a long sequence of accesses to the same cache bank




CacheBleed Operation

- Spy generate a long sequence of accesses to the same cache bank
- Victim accesses to a different cache bank do not affect speed



CacheBleed Operation

- Spy generate a long sequence of accesses to the same cache bank
- Victim accesses to a different cache bank do not affect speed
- Victim accesses to same cache bank cause delays



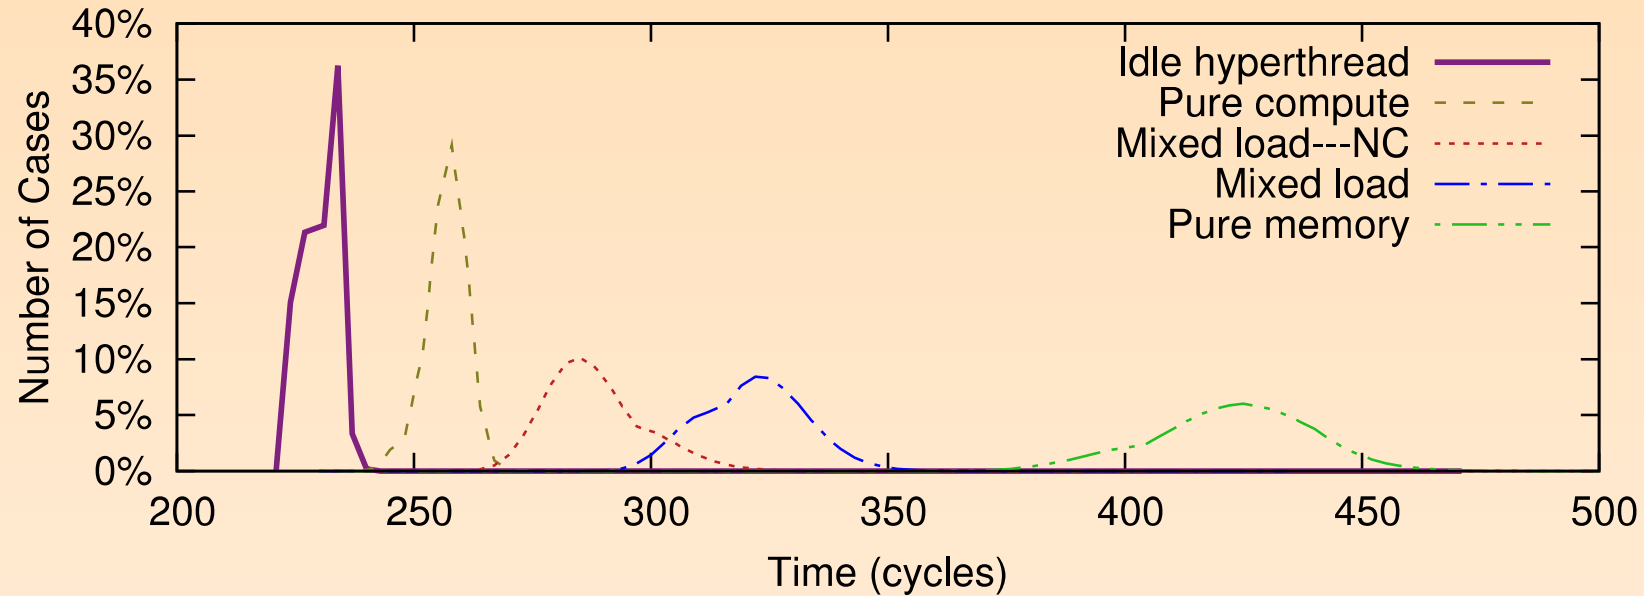
Implementation

```
1      rdtscp
2      movq    %rax, %r10

4      addl    0x000(%r9), %eax
5      addl    0x040(%r9), %ecx
6      addl    0x080(%r9), %edx
7      addl    0x0c0(%r9), %edi
8      addl    0x100(%r9), %eax
9      addl    0x140(%r9), %ecx
10     addl    0x180(%r9), %edx
11     addl    0x1c0(%r9), %edi
.
.
.
256    addl    0xf00(%r9), %eax
257    addl    0xf40(%r9), %ecx
258    addl    0xf80(%r9), %edx
259    addl    0xfc0(%r9), %edi


261    rdtscp
262    subq    %r10, %rax
```

CacheBleed timing



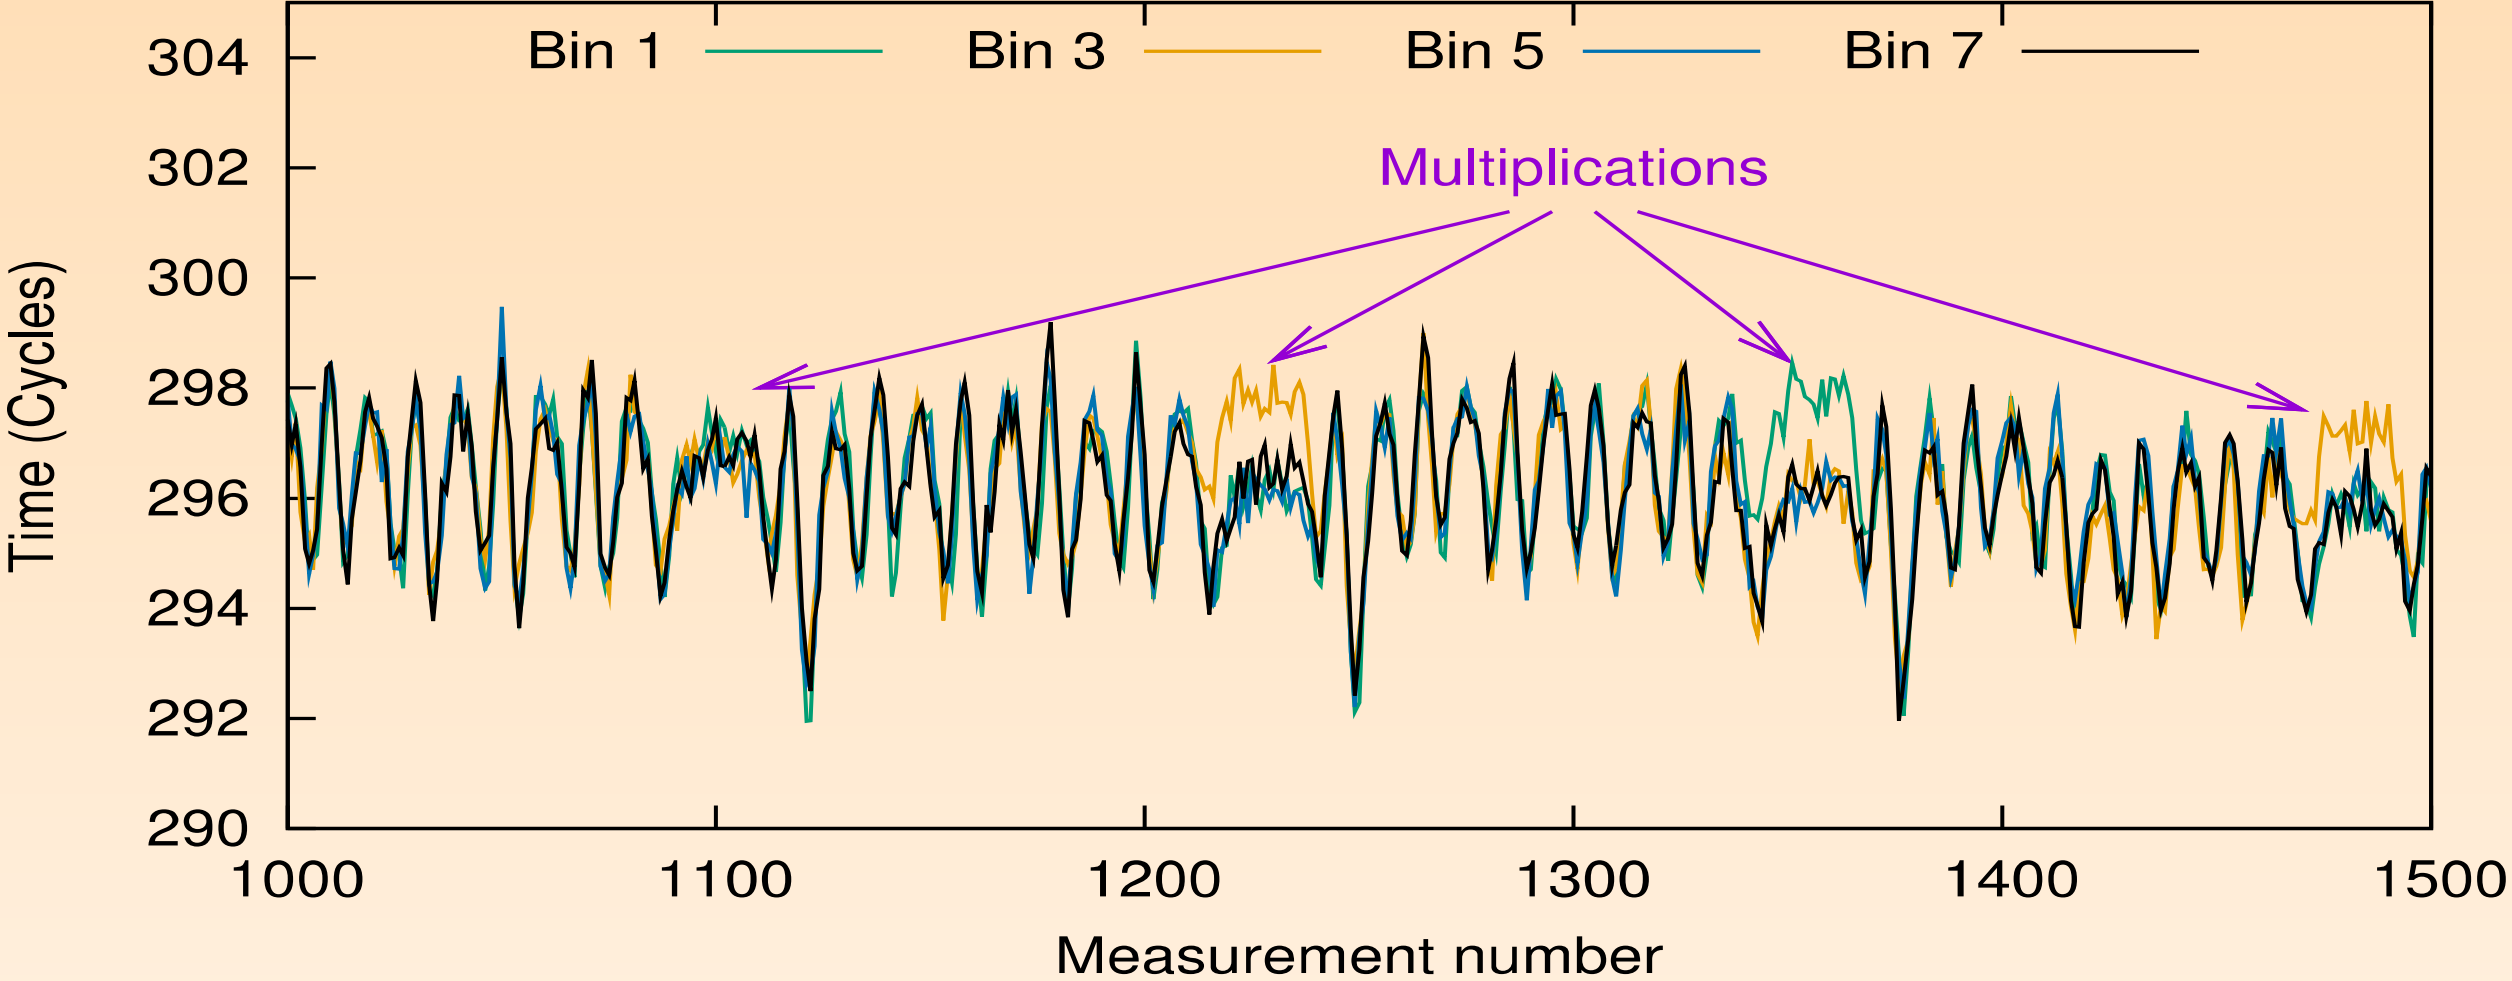
- Need multiple samples to determine cache-bank conflicts

CacheBleed on OpenSSL




- Average of 1,000 sequences on each bin
- Odd and even bins have different timing characteristics


CacheBleed on OpenSSL - Details




Clock Drift



Low-pass filter



Normalised + resampled



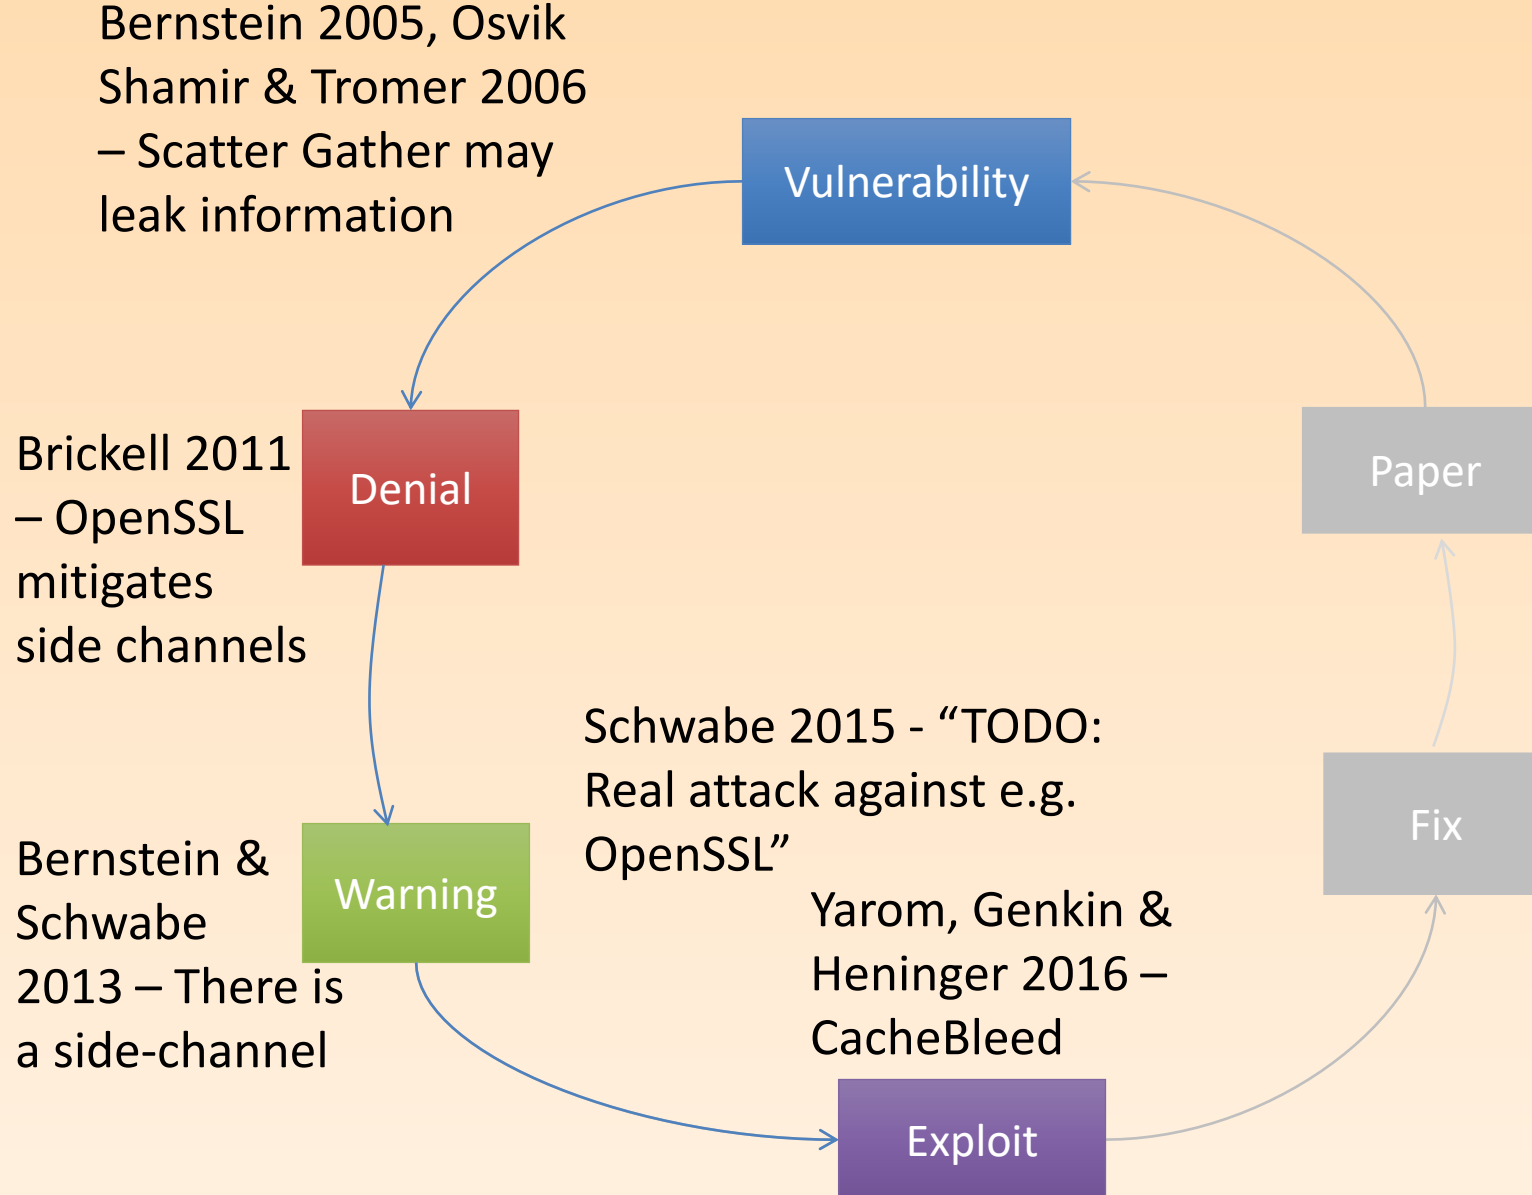
Results

- 16,000 decryptions (1,000 sequences per bin per exponentiation)
 - Less than 5 minutes online attack
- Recover three bits of each multiplier
 - Miss the first and last one or two multipliers

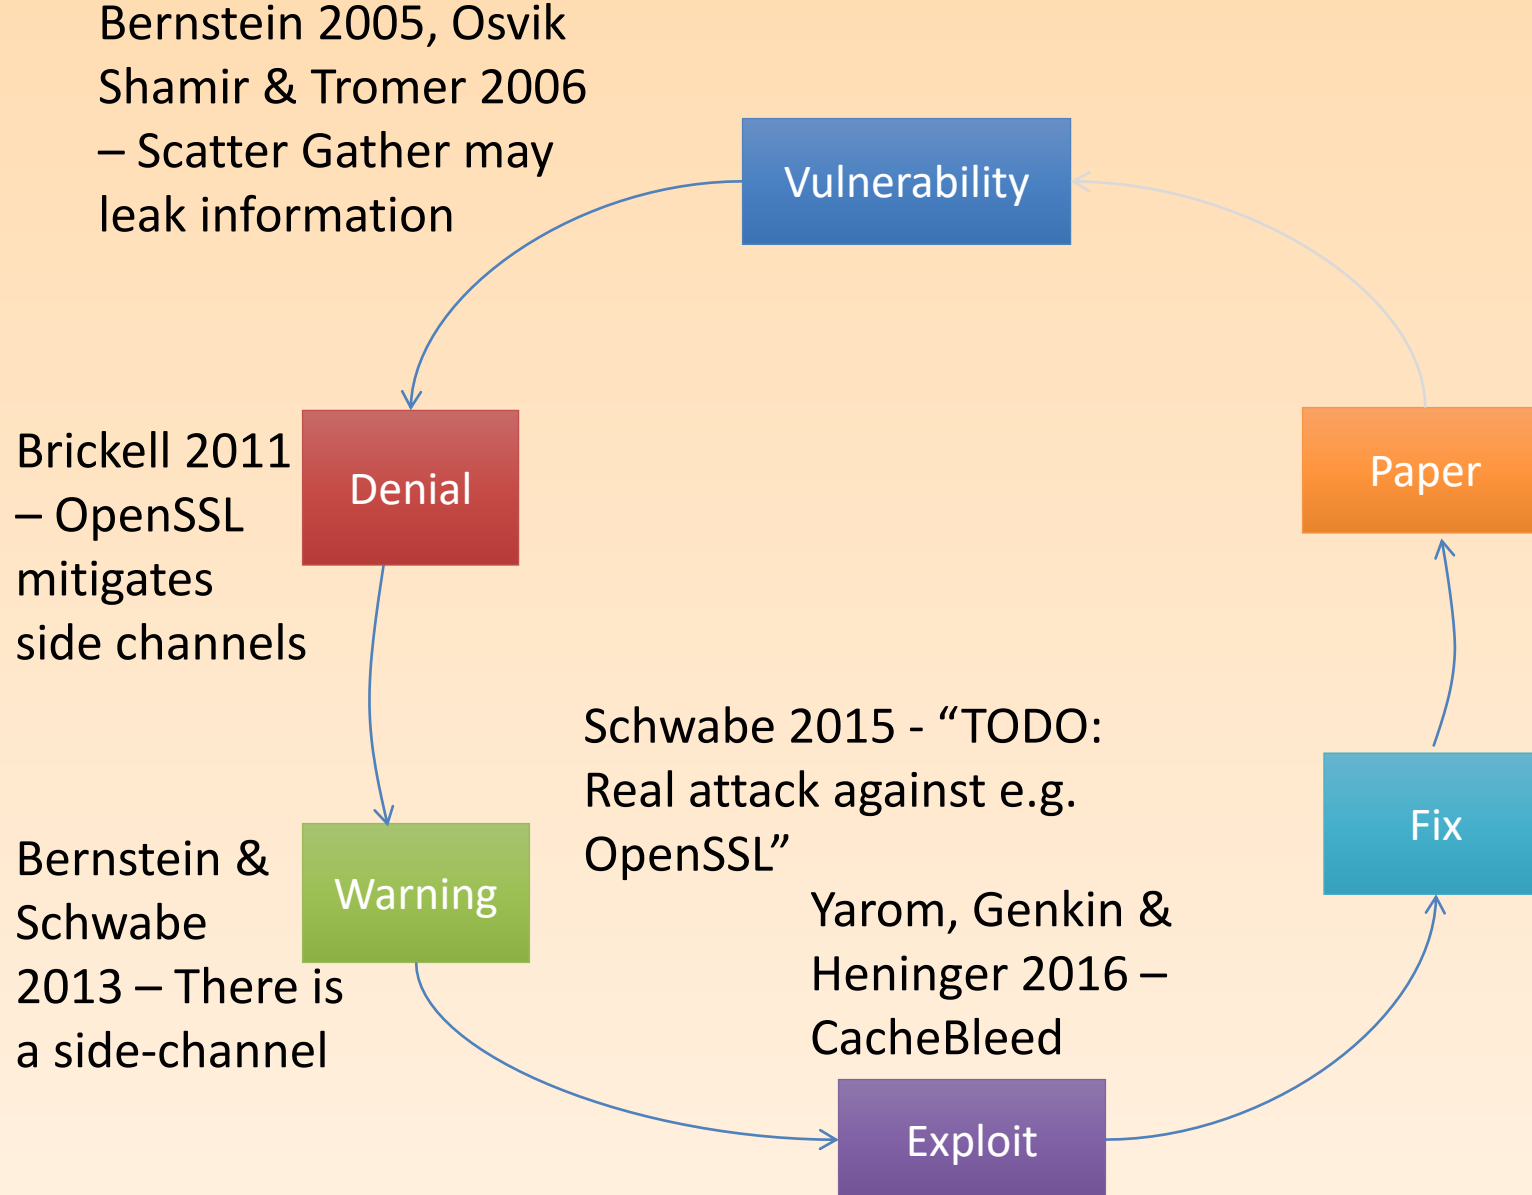
Recovering missing bits

- We know 60% of the bits (3 bits in each 5) in both d_p and d_q
 - Heninger-Shacham requires 50% of the bits
 - Complete key recovery requires two CPU hours – less than 3 minutes on a high-end server

Round 2



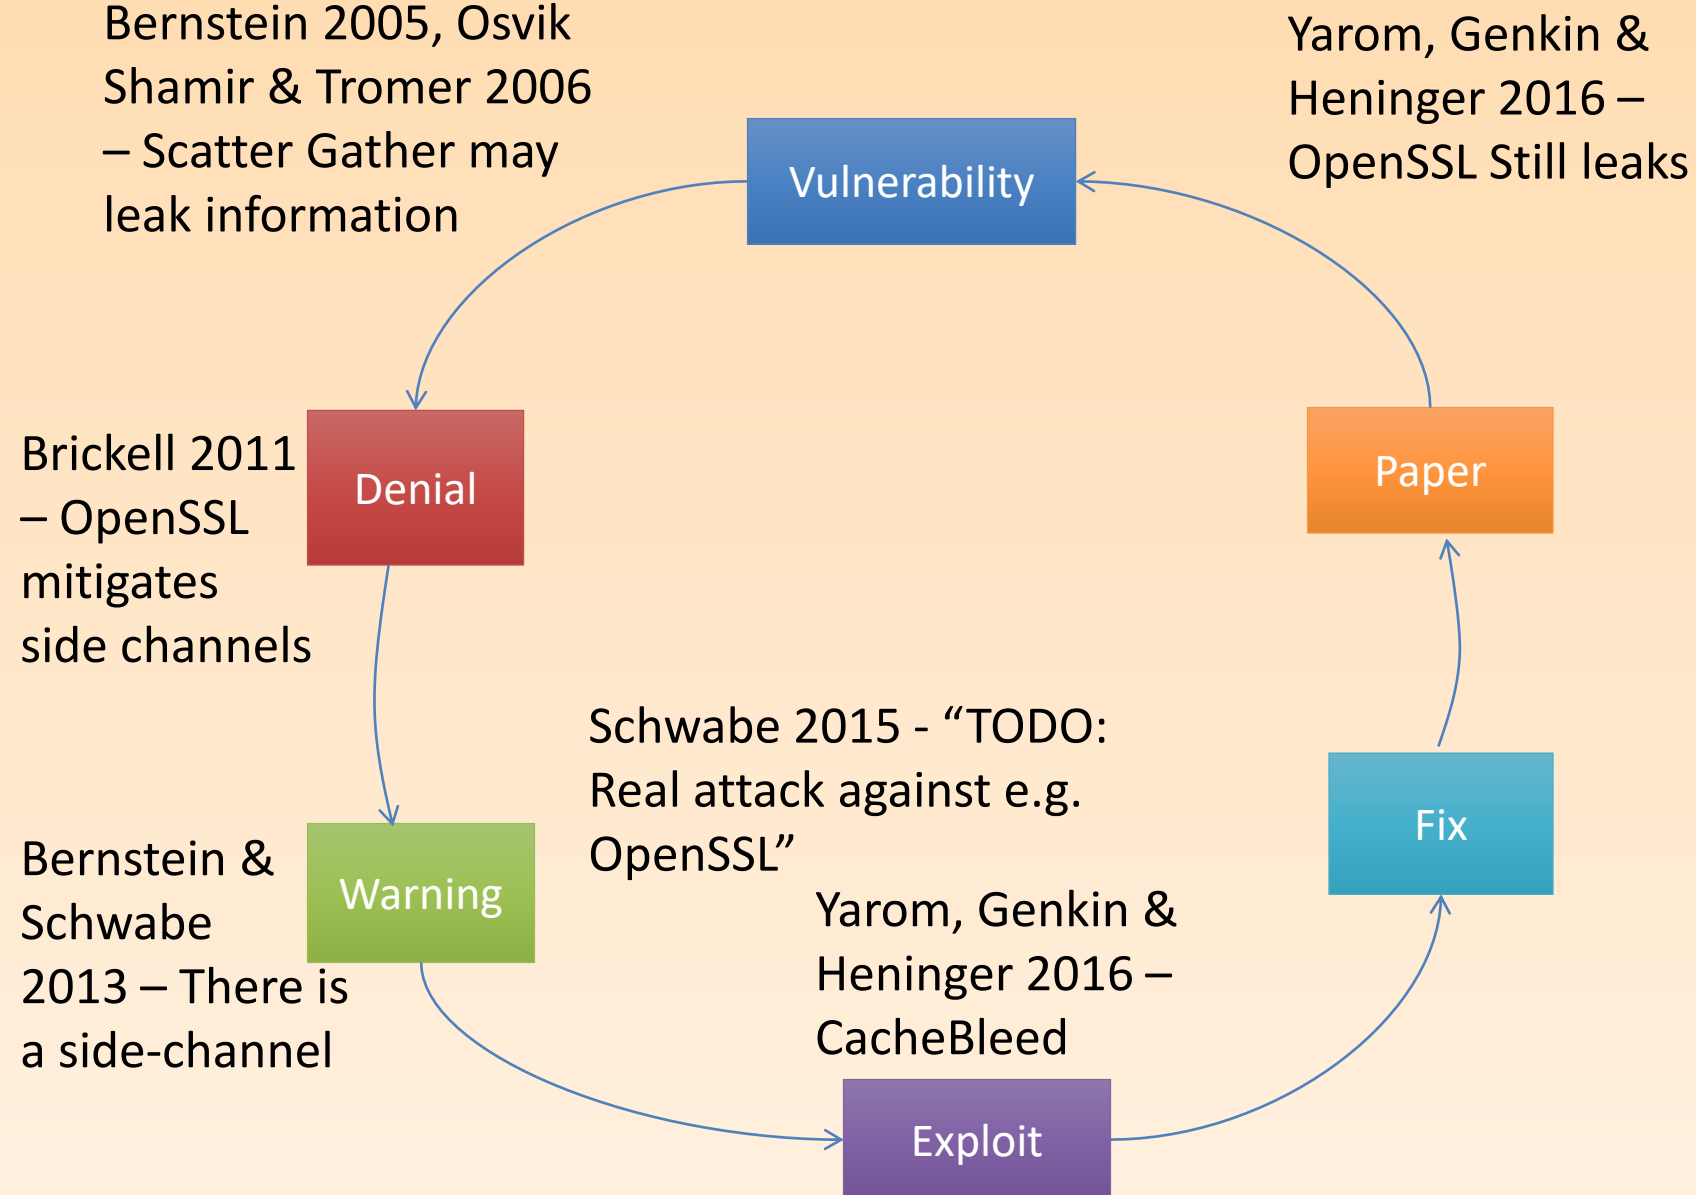
Round 2



OpenSSL Proposed “Fix”

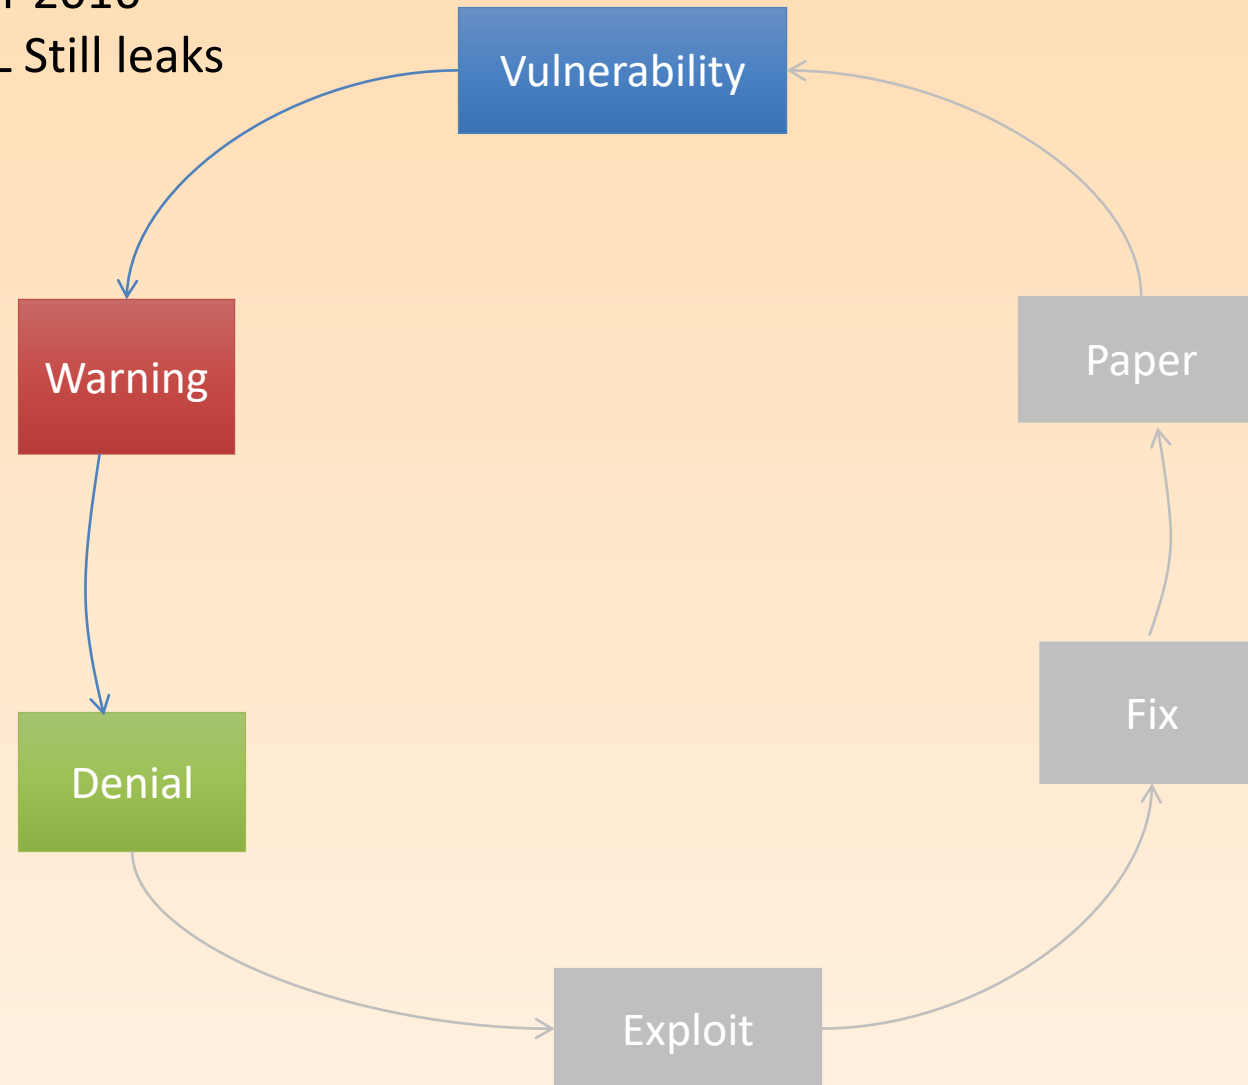
- Use 128-bit reads with masking
 - Only leaks 2 bits per multiplier – not enough for Heninger-Shacham
- Read at a different offset in each of the four cache lines
 - Order depends on the multiplier
 - Too fast for our attack

Round 2



Round 3

Yarom, Genkin &
Heninger 2016 –
OpenSSL Still leaks



CacheBleed

- Fixed in the SkyLake microarchitecture
 - Multiple cache ports
- MemJam – false dependencies
 - Moghimi et al. “MemJam: A False Dependency Attack against Constant-Time Crypto Implementations”, CT-RSA 2018
- Port contention
 - Aldaya et al. “Port Contention for Fun and Profit”, IEEE SP 2019

Summary

- Microarchitectural attacks often return partial information
 - Can use redundancy to reconstruct key
 - Heninger-Shacham algorithm
-
- Next: lower-level caches and eviction sets
 - Read: Vila et al. “Theory and Practice of Finding Eviction Sets”, IEEE S&P 2019